

# NoSQL in a glimpse

[https://github.com/federicodassereto/MAIA\\_NoSQL](https://github.com/federicodassereto/MAIA_NoSQL)

Federico Dassereto\*  
federico.dassereto@edu.unige.it

Dibris

\*Giovanna Guerrini  
\*Dan Suciu



UNIVERSITÀ  
DEGLI STUDI  
DI GENOVA

Not every data management/analysis problem is best solved using a traditional DBMS.

Database Management System (DBMS) provides efficient, reliable, convenient, and safe multi-user storage of and access to massive amounts of persistent data.

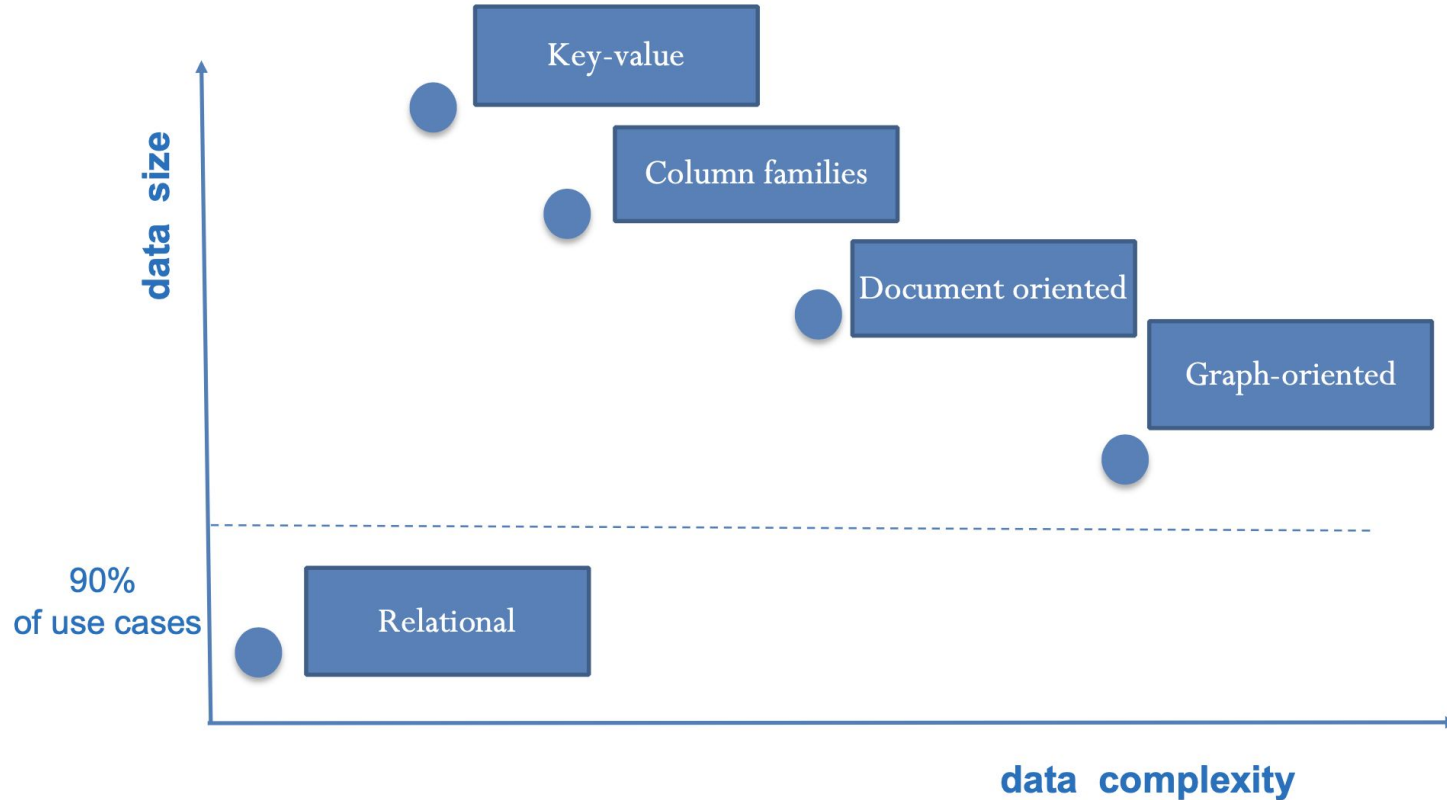
## Alternative to traditional relational DBMS

- Flexible schema
- Quicker/cheaper to set up
- Massive scalability
- Relaxed consistency higher performance & availability

## Drawbacks

- No declarative query language => more programming
- Relaxed consistency => fewer guarantees

# NoSQL data models



## Focus on

Document Oriented

CouchDB

**MongoDB**

SimpleDB

Graph Oriented

Neo4J

**Gremlin**

.....



## Document-oriented Data Model

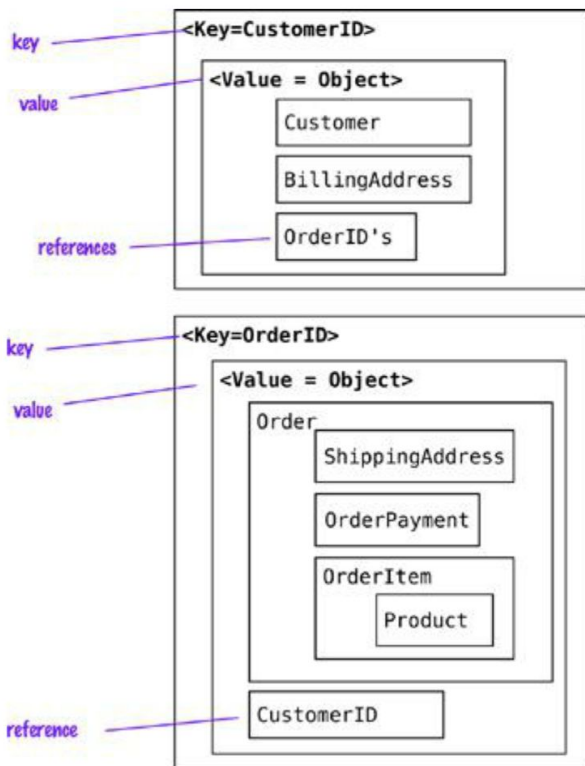
Each data instance is represented in the form (key, document)

- key is an identifier
- document is the aggregate, corresponding to a set of <name, nested-document> pairs

Usually nested-documents are represented according to a semi-structured data model like XML, RDF, JSON

Structure of the aggregate visible

# Document-oriented Data Model



# Customer object

```
{
  "1": {
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}]
  }
}
```

# Order object

```
{
  "99": {
    "order": {
      "orderDate": "Nov-20-2011",
      "orderItems": [{"productId": 27, "price": 32.45}],
      "orderPayment": [{"ccinfo": "1000-1000-1000-1000",
        "txnId": "abelif879rft"}],
      "shippingAddress": {"city": "Chicago"}
    }
  },
  "customerId": 1
}
```

Sub-document



## Document-oriented - Interaction

Basic lookup based on the key and field names in the aggregates

- `Void put(key, document)`
- `Document get(key)`
- `Void set(key, name, value)`
- `Document get(key, name)`
- `Void remove(key)`

# Example

We can query inside  
(nested) documents, at  
any level:

- find all orders issued in  
November 2011  
(condition on OrderDate)
- find orderPayment  
information associated  
with order 99

```
# Customer object
{
  "1": {
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}]
  }
}


# Order object
{
  "99": { "order": {
    "orderDate": "Nov-20-2011",
    "orderItems": [{"productId": 27, "price": 32.45}],
    "orderPayment": [{"ccinfo": "1000-1000-1000-1000",
                      "txnId": "abelif879rft"}],
    "shippingAddress": {"city": "Chicago"}
  }
  "customerId": 1
}
```

# Graph data stores

- Graph databases are motivated by a different frustration with relational databases
  - Complex relationships require complex join
  - Modeling of many-to-many relationships
- Goal
  - Capture data consisting of complex relationships
  - Data naturally modelled as graphs
    - Social graphs
    - Web & Semantic Webs
    - Networks (roads, rails, ...)
- A graph is a generic, intuitive, and well-known data structure

# Graph data stores: Vertices and Edges

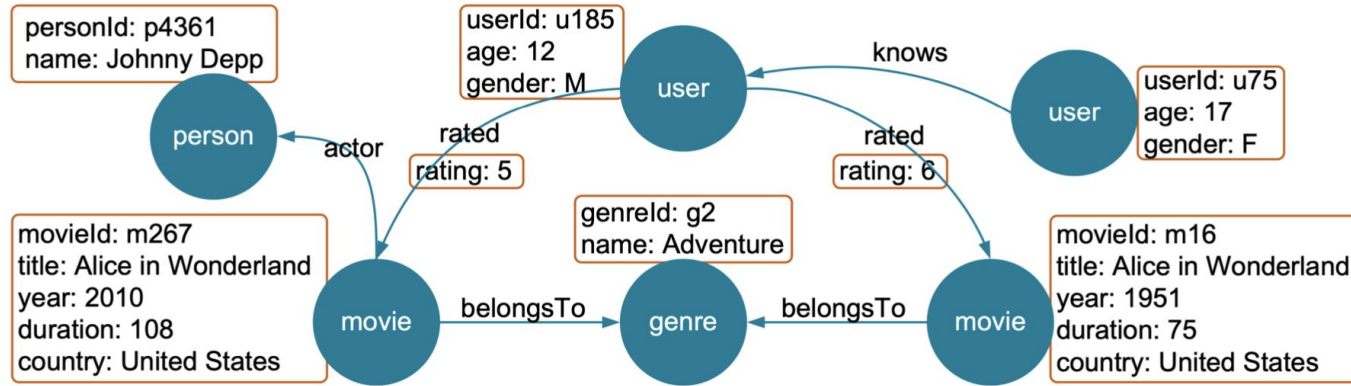
- A graph  $G=(V,E)$  is a pair of vertices and edges connecting them



Vertices (or nodes)  
– Label: E.g. person, movie

Edges  
– Type: E.g. actor, rated, belongsTo

# Properties



- Vertices (or nodes)
  - Label: E.g. person, movie
  - Properties: E.g. name, age
- Edges
  - Type: E.g. actor, rated, belongsTo
  - Properties: E.g. rating

# Schema less data model

- Each entity (object) can have different properties, just like a document database
- Any entity can have a relationship with any other entity
- Relationships have a type, and any pair of entities can have a relationship of any type
- Relationships have properties, so can be thought of as entities that join other entities
- Entity pairs can have more than one relationship

# Exercises



mongoDB

Harry Potter

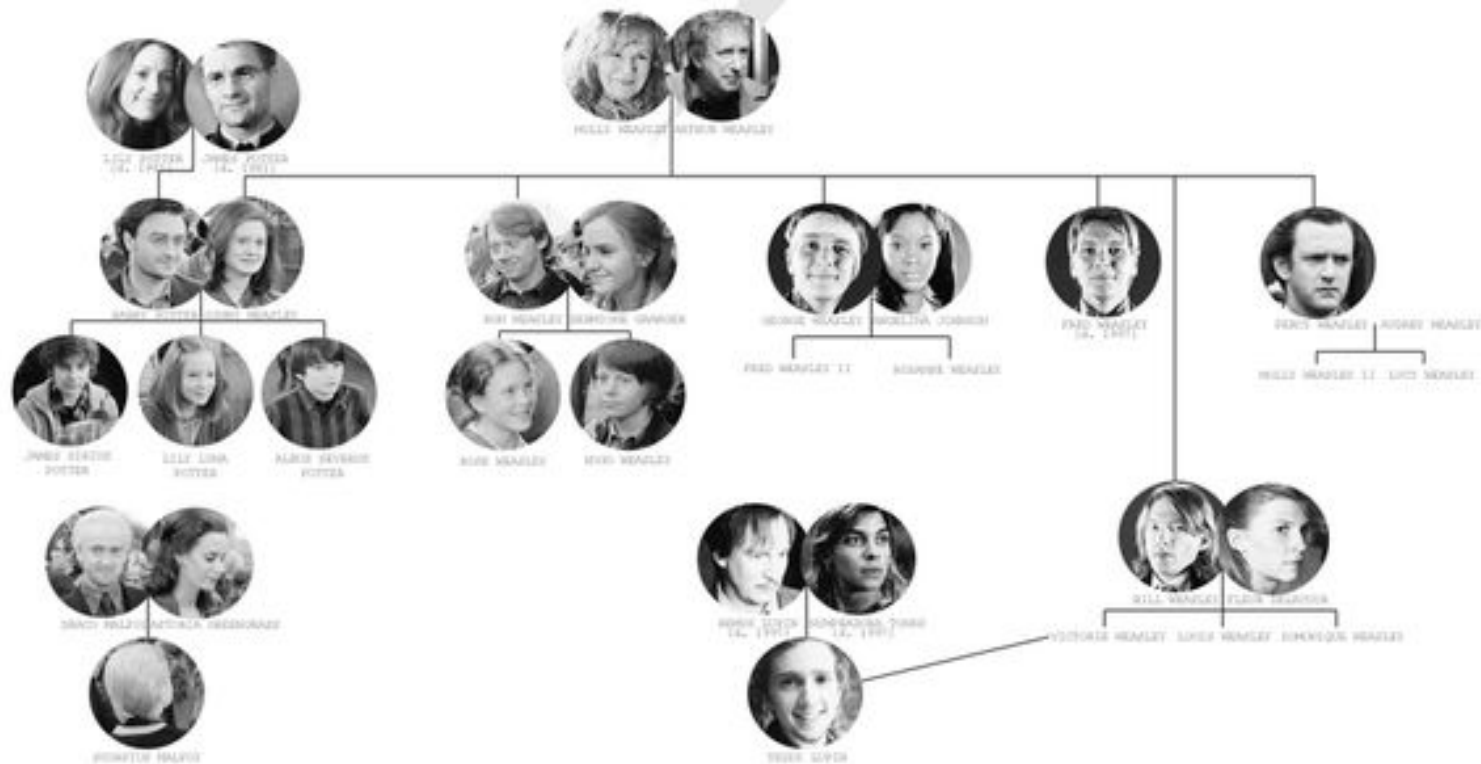


**Gremlin**  
 $G = (V, E)$

how i  
met your  
mother

# Harry Potter

## FAMILY TREES





Why mapping it via documents instead of graph?

## Workflow

Create a document for each family (at least 3)

Query the DB to obtain simple informations

Which MongoDB function is the best one?

Try to add new unrelated documents

Try to retrieve documents by common features

Is it easier than joins on SQL?

How I met your mother  
relationships among characters

## Workflow

Create a graph  $G=(V,E)$  where  $V$  are characters and  $E$  are relations

Query the DB to obtain simple informations

How do we traverse the graph?

Try to filter nodes by 'Location'

Try to add new Nodes and Relations

Try to simulate an SQL join

Is it easier than joins on SQL?